

# FRAME: FAULT-TOLERANT REAL-TIME MESSAGING ARCHITECTURE

**Paper By: Chao Wang, Christopher Gill, Chenyang Lu**

**Presented By: Gregor Peach**

WHAT IS FRAME?

# BASIC PROBLEM

- We have IoT devices
- They need to communicate
- Message passing!!!
- Goals
  - Don't lose messages
  - Messages go
    - zoom
    - Zoom




Fig. 1. An Illustration of IIoT Edge Computing.

# A SYSTEM HAS MANY PARTS


## Heterogeneous requirements

- Emergency Stop Mechanism
  - No message loss
  - 10ms latency needed
- Logger
  - No message loss
  - 1sec latency okay
- Dashboard Updater
  - Occasional message loss
  - 100ms latency desired




# MESSAGING DESIGN

- Two Basic Designs of Messaging
  - Brokered
  - Unbrokered
  - Q: What are the pros/cons of each?
- Broker Placement
  - Cloud
  - Edge
  - Q: Which is better for IoT?




# TOPICS

- Broker handles many message streams
- A Topic is like a email-list for messages
- Topic operations
  - Send/Publish on topic
  - Receive/Subscribe to topic
- Q: How is this different from a set of queues?



# FRAME IS MIDDLEWARE

- Software for sending messages to topics
- Software for subscribing to topics
  - Also has to handle asynchronously receiving messages
- A bunch of servers serving as message brokers



# FRAME (THE NOVEL PARTS)

- Each Topic has metadata
  - What many messages in a row can be lost?
  - What is the deadline for messages on this channel?  
(Real time!!!)
- FRAME Guarantees
  - Always meet loss prevention promises
  - If there are no faults
    - Always meet message deadlines
  - If there are faults
    - Recovery is fast so deadlines aren't too badly compromised



# CONTRIBUTION BREAKDOWN

# MODEL FOR FAULT TOLERANT REAL TIME MESSAGING

“A new fault-tolerant real-time messaging model. We describe timing semantics for message delivery, identify under what conditions a message may be lost, prove timing bounds for real-time fault-tolerant actions in terms of traffic/service parameters, and demonstrate how the timing bounds can support efficient and appropriate message differentiation to meet each requirement.” (FRAME pg 2)

# A SYSTEM BUILT AROUND THAT MODEL

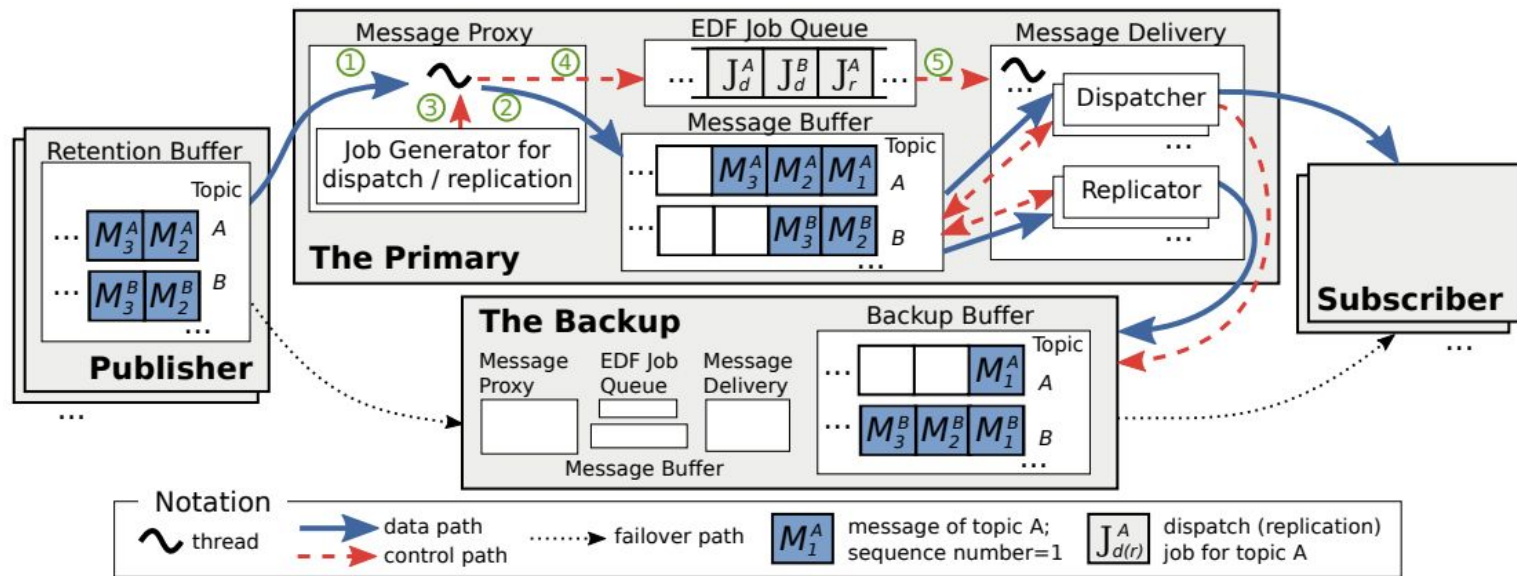


Fig. 4. The FRAME Architecture.

# A SYSTEM BUILT AROUND THAT MODEL

- Avoiding Data Loss
  - There are two servers, a primary and a backup
  - The backup becomes the primary (and visa versa) if the primary crashes
    - Messages are only replicated if needed for guarantees
  - Q: Can messages be delivered twice?

TABLE 3  
ALGORITHM FOR DISPATCH-REPLICATE COORDINATION.

Type of Operation	Procedure
Dispatch	<ol style="list-style-type: none"> <li>1. dispatch the message to the subscriber</li> <li>2. set <i>Dispatched</i> to True</li> <li>3. if <i>Replicated</i> is True, request the Backup to set <i>Discard</i> to True</li> </ol>
Replicate	<ol style="list-style-type: none"> <li>1. if <i>Dispatched</i> is True, abort</li> <li>2. replicate the message to the Backup</li> <li>3. set <i>Replicated</i> to True</li> </ol>
Recovery (in the Backup)	<ol style="list-style-type: none"> <li>1. if <i>Discard</i> is True, skip the message</li> <li>2. create a dispatching job for the message</li> <li>3. push the job into the EDF Job Queue</li> </ol>

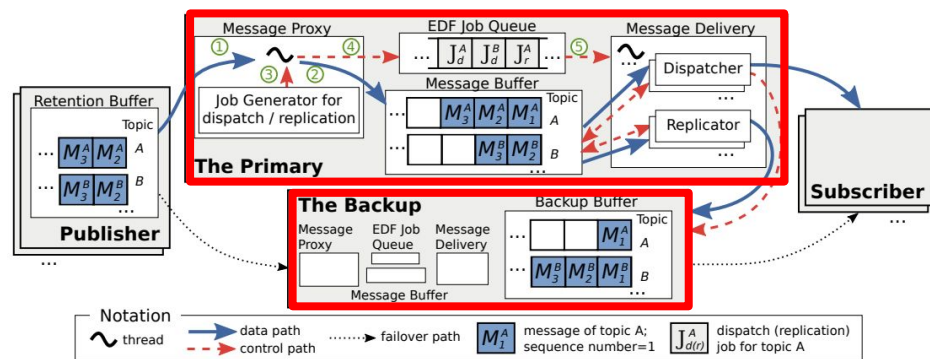


Fig. 4. The FRAME Architecture.

# A SYSTEM BUILT AROUND THAT MODEL

## - Fast Message Delivery

- Broker lives in the edge (even though it supports cloud clients)
- Earliest deadline first message transmission
- Q: Why EDF? What are the pros and cons?

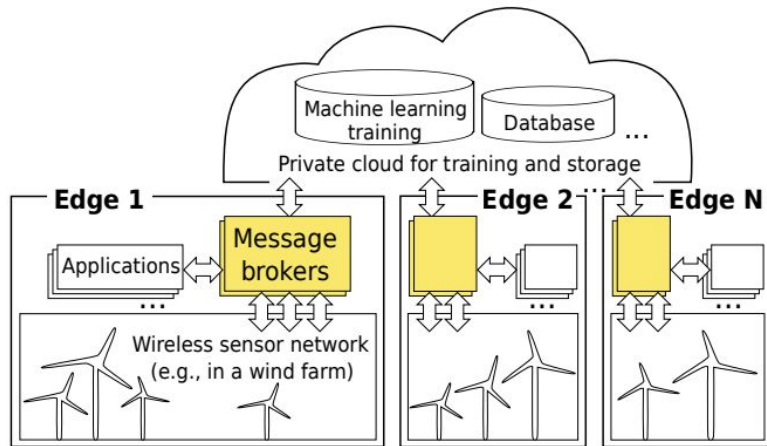
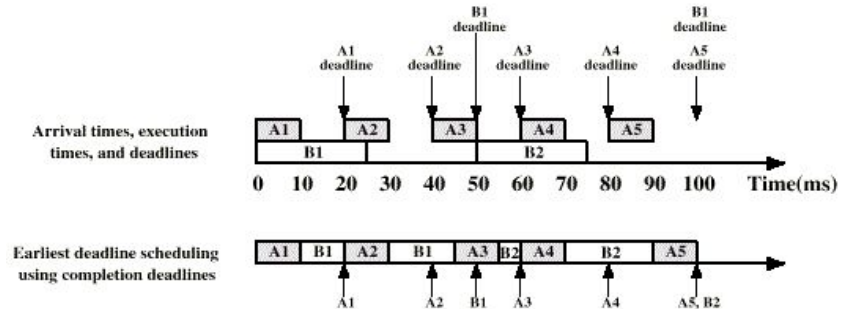


Fig. 1. An Illustration of IIoT Edge Computing.



EVALUATION

# LOSS PREVENTION SUCCESS RATES

- The Candidates
  - FRAME
    - I just tried to explain this
  - FRAME+
    - FRAME + publisher message retention
  - FCFS (first-come first-serve)
    - Delivers messages as they come
    - Coordinates with backup whenever a message is sent (Running on same basic framework as FRAME)
  - FCFS- (first-come first-serve-)
    - Very limited coordination with backup
    - Only occasionally lets backup know what it has sent

# LOSS TOLERANCE SUCCESS RATE

$D_i$  = deadline (in ms)

$L_i$  = max allowable

Losses in a row

TABLE 4  
SUCCESS RATE FOR LOSS-TOLERANCE REQUIREMENT (%).

$D_i$	$L_i$	FRAME+	FRAME	FCFS	FCFS-
Workload = 13525 Topics					
50	0	100.0	$80.0 \pm 30.1$	0.0	100.0
50	3	100.0	$80.0 \pm 30.1$	0.0	100.0
100	0	100.0	$73.2 \pm 30.7$	0.0	$78.4 \pm 13.3$
100	3	100.0	$79.3 \pm 29.9$	0.0	$99.3 \pm 0.5$
100	$\infty$	100.0	100.0	100.0	100.0
500	0	100.0	$80.0 \pm 30.1$	0.0	100.0

Note: Test includes fault injection



# LOSS PREVENTION SUCCESS RATES: ANALYSIS




- FCFS with backup coordination sucks
  - Spends too much time backing stuff up
- FCFS- saves some of that time
  - Doesn't coordinate with backup
  - Backup will save messages already sent
  - So recovering from the backup is slow (even though there is little data loss)
- FRAME
  - Is smart about coordinating with backup
    - Doesn't backup every message
- FRAME+
  - Shows the power of publisher retention

TABLE 4  
SUCCESS RATE FOR LOSS-TOLERANCE REQUIREMENT (%).




$D_i$	$L_i$	FRAME+	FRAME	FCFS	FCFS-
Workload = 7525 Topics					
50	0	100.0	100.0	0.0	100.0
50	3	100.0	100.0	0.0	100.0
100	0	100.0	100.0	0.0	100.0
100	3	100.0	100.0	0.0	100.0
100	$\infty$	100.0	100.0	100.0	100.0
500	0	100.0	100.0	0.0	100.0
Workload = 10525 Topics					
50	0	100.0	100.0	0.0	100.0
50	3	100.0	100.0	0.0	100.0
100	0	100.0	100.0	0.0	100.0
100	3	100.0	100.0	0.0	100.0
100	$\infty$	100.0	100.0	100.0	100.0
500	0	100.0	100.0	0.0	100.0
Workload = 13525 Topics					
50	0	100.0	80.0 $\pm$ 30.1	0.0	100.0
50	3	100.0	80.0 $\pm$ 30.1	0.0	100.0
100	0	100.0	73.2 $\pm$ 30.7	0.0	78.4 $\pm$ 13.3
100	3	100.0	79.3 $\pm$ 29.9	0.0	99.3 $\pm$ 0.5
100	$\infty$	100.0	100.0	100.0	100.0
500	0	100.0	80.0 $\pm$ 30.1	0.0	100.0

# FAULT RECOVERY PERFORMANCE




FRAME+ =>






FRAME =>



FCFS =>



FCFS- =>



(a) Category 0,  $T_i = 50$ ,  $D_i = 50$ .

(b) Category 2,  $T_i = 100$ ,  $D_i = 100$ .

(c) Category 5,  $T_i = 500$ ,  $D_i = 500$ .

HEAVY LOAD

MEDIUM LOAD

LIGHT LOAD

# REAL-TIME PERFORMANCE

TABLE 5  
SUCCESS RATE FOR LATENCY REQUIREMENT (%).

		Workload = 13525 Topics			
50	0	$98.4 \pm 2.9$	$85.4 \pm 21.7$	$0.1 \pm 0.1$	$99.4 \pm 3.6\text{E-}1$
50	3	$98.4 \pm 2.9$	$85.3 \pm 21.7$	$0.2 \pm 0.2$	$99.5 \pm 2.3\text{E-}1$
100	0	$97.6 \pm 4.4$	$83.7 \pm 21.9$	$2.6\text{E-}4 \pm 6.0\text{E-}4$	$98.3 \pm 1.0$
100	3	$97.6 \pm 4.4$	$83.8 \pm 21.9$	$9.9\text{E-}4 \pm 2.2\text{E-}3$	$98.3 \pm 1.1$
100	$\infty$	$97.6 \pm 4.4$	$83.8 \pm 21.9$	$6.6\text{E-}4 \pm 1.5\text{E-}3$	$98.3 \pm 1.1$
500	0	$98.6 \pm 2.8$	$86.1 \pm 21.8$	0.0	100.0

CRITIQUES

# FROM THE GITHUB ISSUE...

@hjaensch7: Some figures seemed inefficient?

@RyanFisk2: What happens if the backup fails? Or if both fail?

@tuhinadasgupta: What are the security concerns of FRAME?