

CHAOS


Paper By: PK Gadepalli, G Peach, G Parmer, J Espy, Z Day
Presentation By: Gregor Peach

Motivation

Mixed Criticality Systems

- Hear me out, but what if have a system that does

TWO THINGS!!!!
(big wow)



Low and High Criticality

High Criticality

- Important tasks
- Highly tested code



Low Criticality

- Unimportant tasks
- Less tested code



State of the Art: Real-time Systems

Classic Real-time Operating Systems


- eg. FreeRTOS
- Pros:
 - Predictable execution (real-time)
 - Single core
 - High assurance
- Cons:
 - No multicore
 - No rich POSIX APIs



State of the Art: POSIX

LINUX

- eg. Ubuntu
- Pros:
 - Rich programming environment
 - Good hardware support
 - Real-time patch available
- Cons:
 - Really big and complex
 - Real-time performance terrible
 - Big surface area for attacks



Mixed Criticality Software Lacks a Clear Home

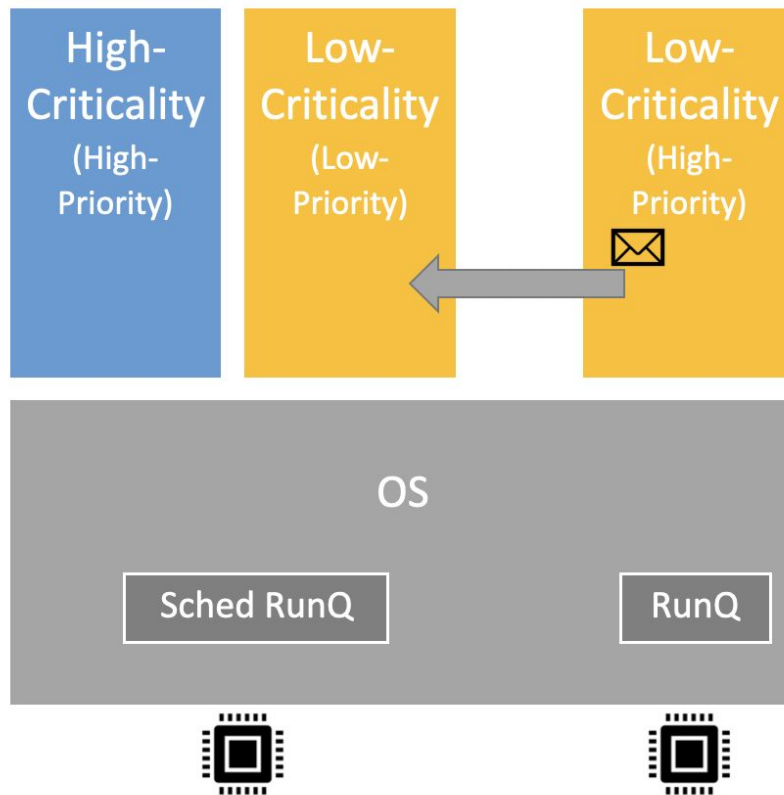
If I have a mixed criticality system, what platform should I build it on?

What are the pros and cons of running it on FreeRTOS?

Linux?




This Setup is Unlikely to Work

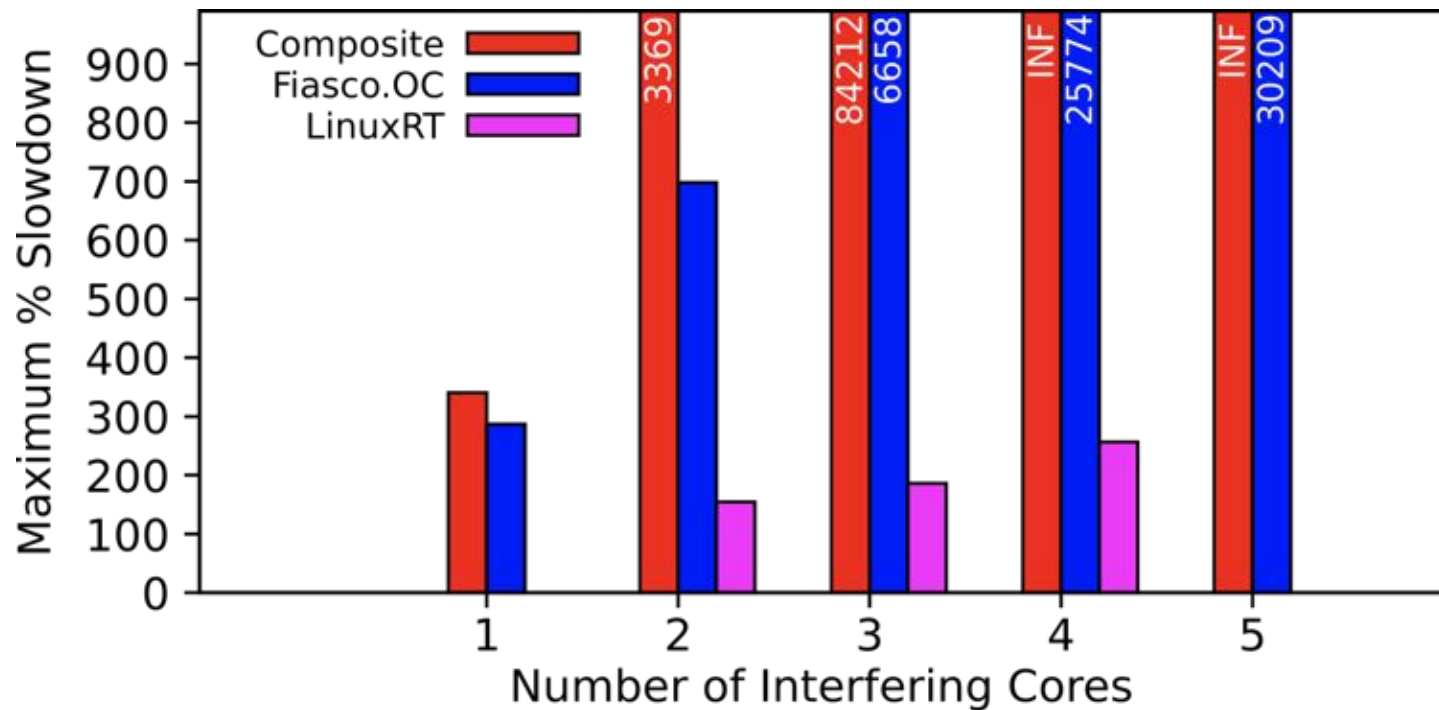


And Even If It Did...

- Naive core sharing has a problem
 - Cross core interrupt denial of service attack

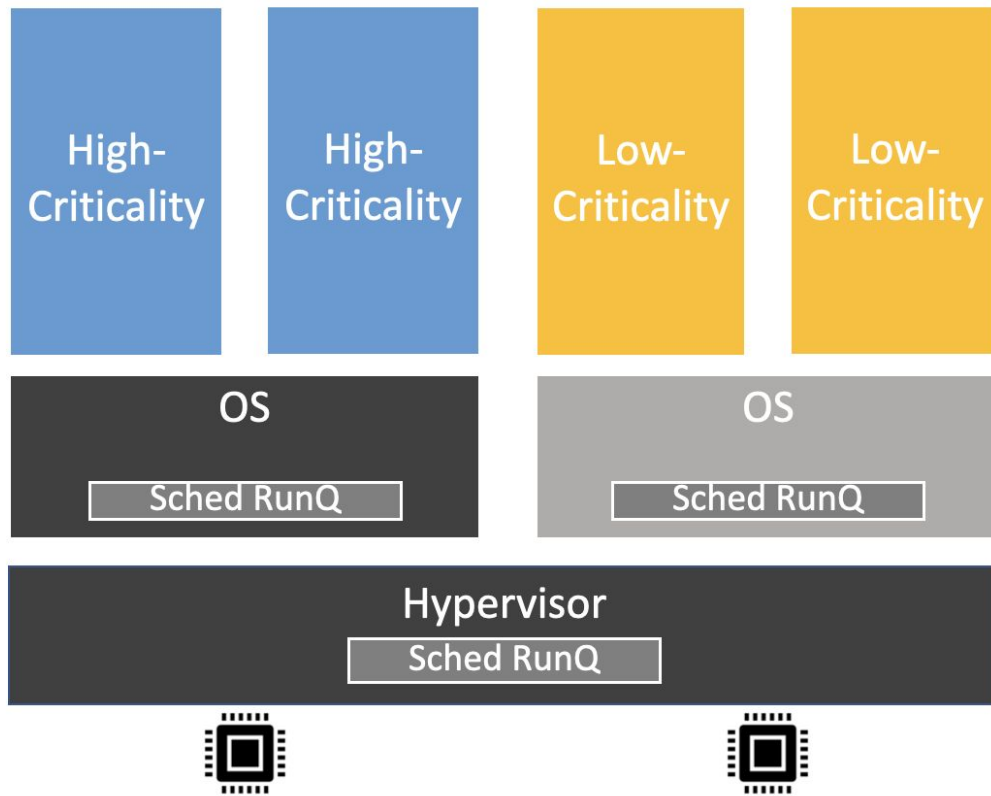


Performance Problem!



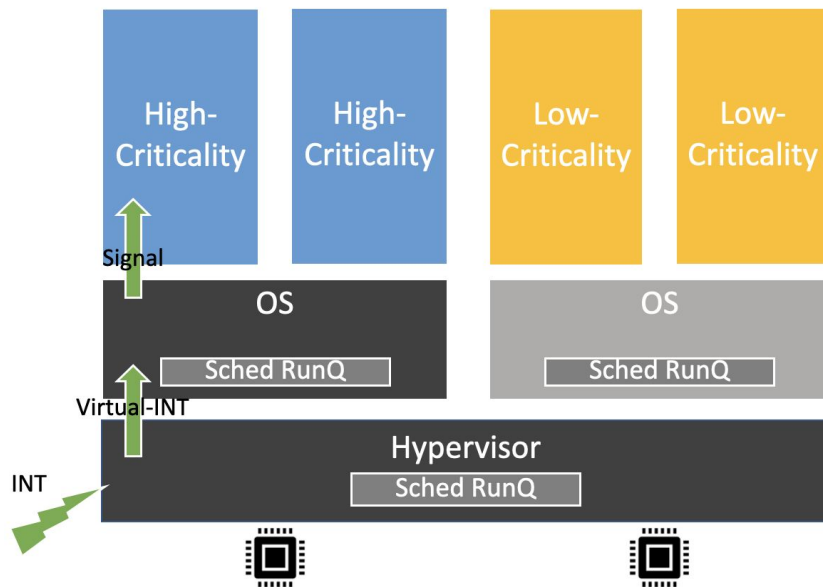
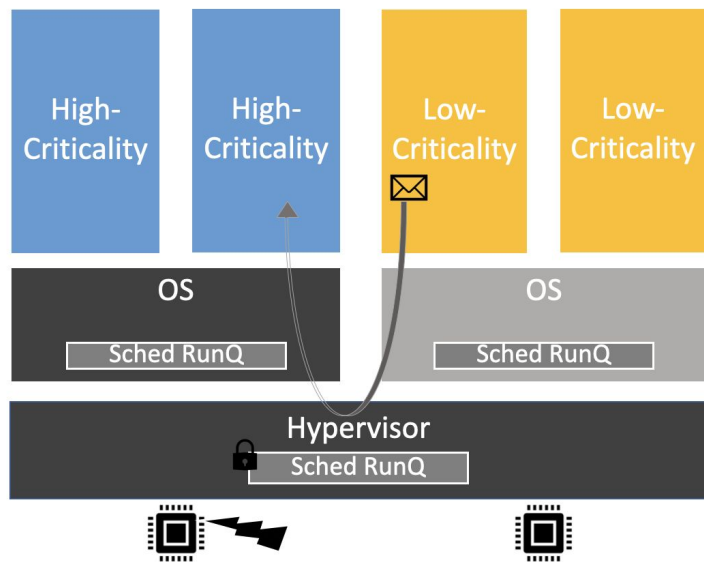
Solution in Practice: Per Core Virtualization

— — —



Solution in Practice: Per Core Virtualization

Interrupts handled properly...

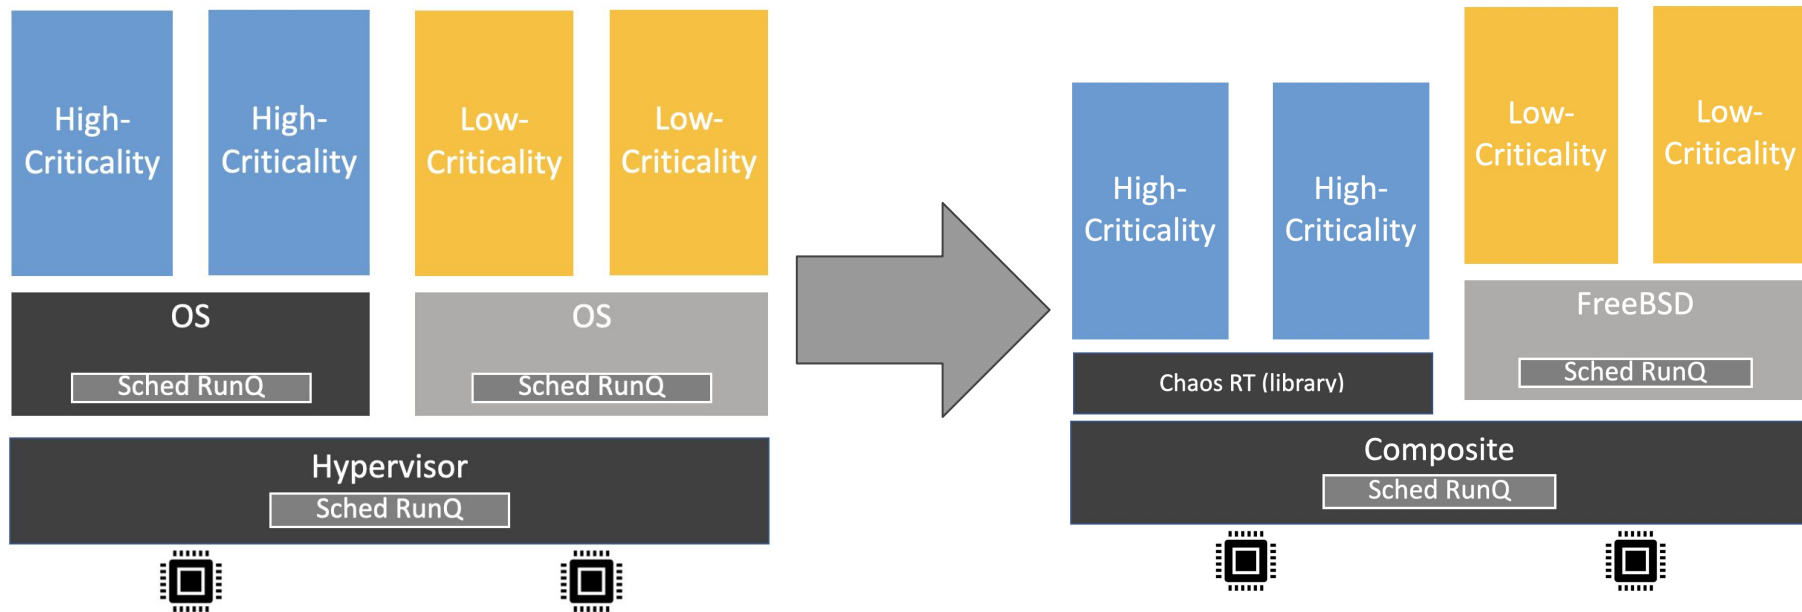


Problems with Per Core Virtualization

- It's super slow, big overheads for normal computation
 - Real-time worst case bounds are subject to variance
(Hierarchical Overheads)
- Slower to coordinate between high and low assurance

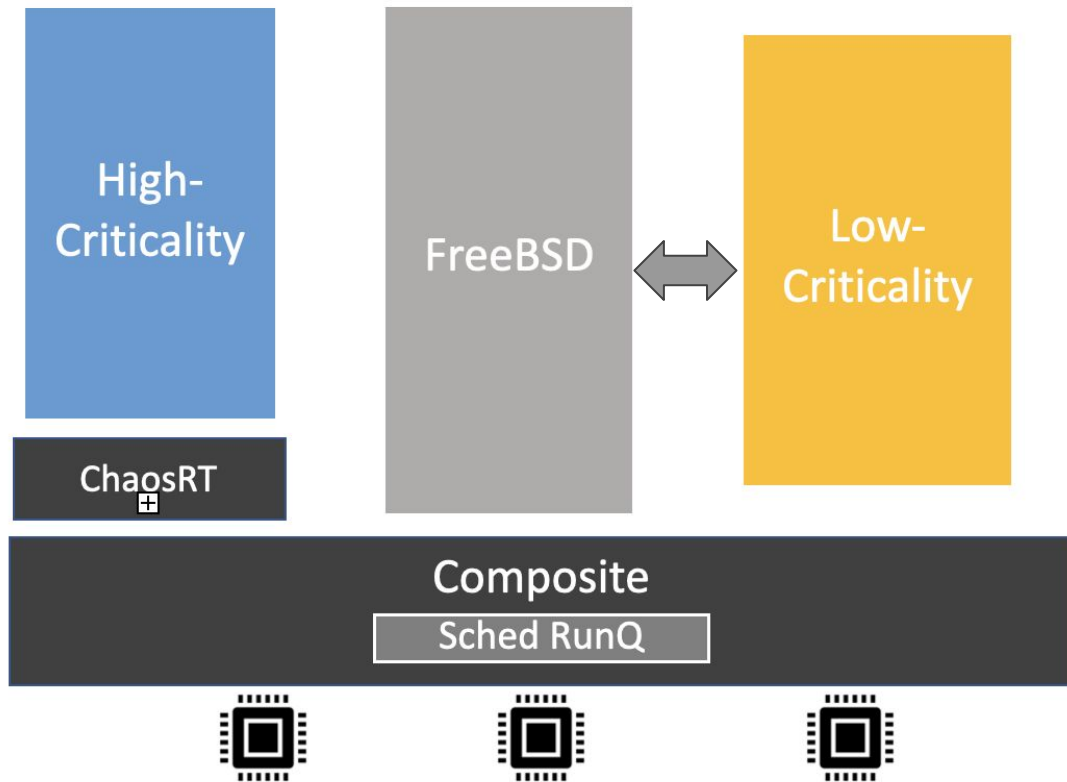
What is CHAOS?

CHAOS: Presenting Devirtualization!!



CHAOS: Presenting Devirtualization!!

— — —

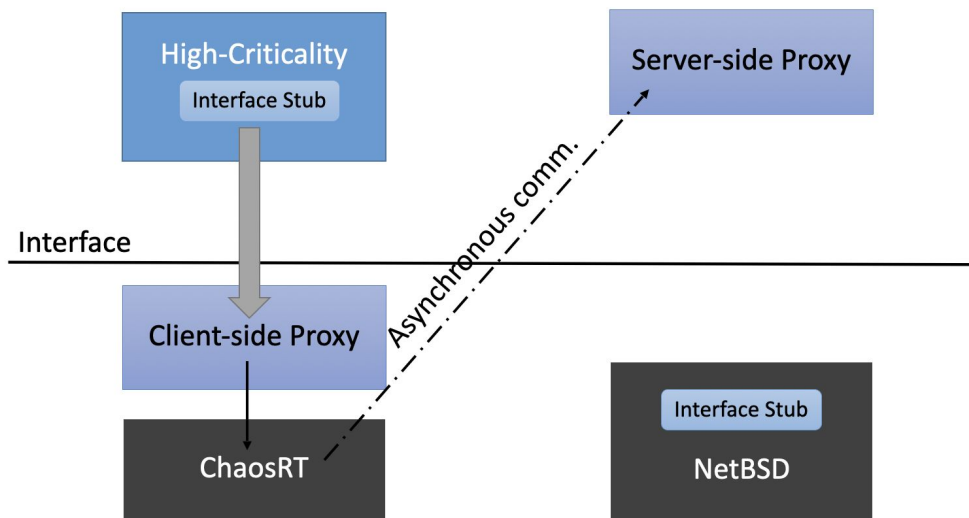


Devirtualization: Get Rid of the Slow VM Layer

- Per Core Virtualization Slow
 - Too many layers of isolation
- Removing it is a win for predictability and performance!
- Works for both low and high criticality code!

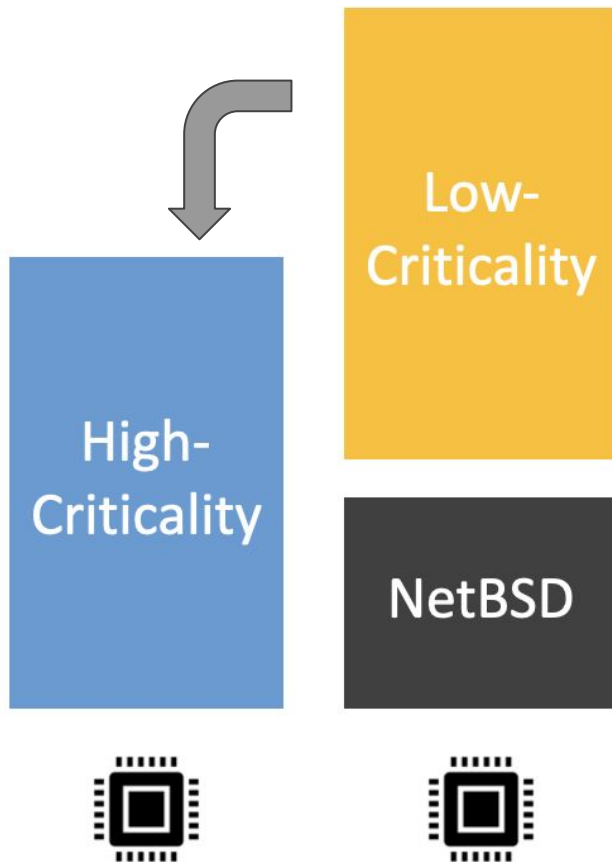
Extension: High Assurance Code uses UNIX APIs

Proxies = mechanism for letting High-Criticality and Low-Criticality processes communicate

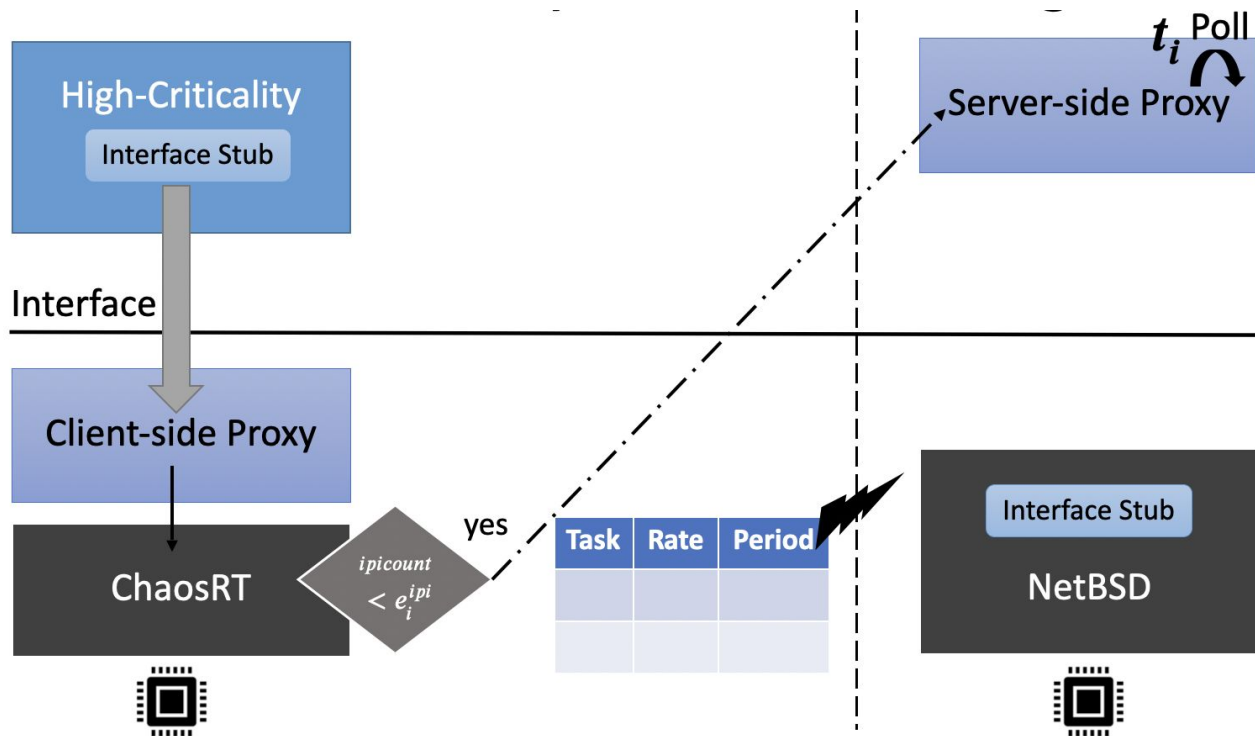


Why Not Talk Directly?

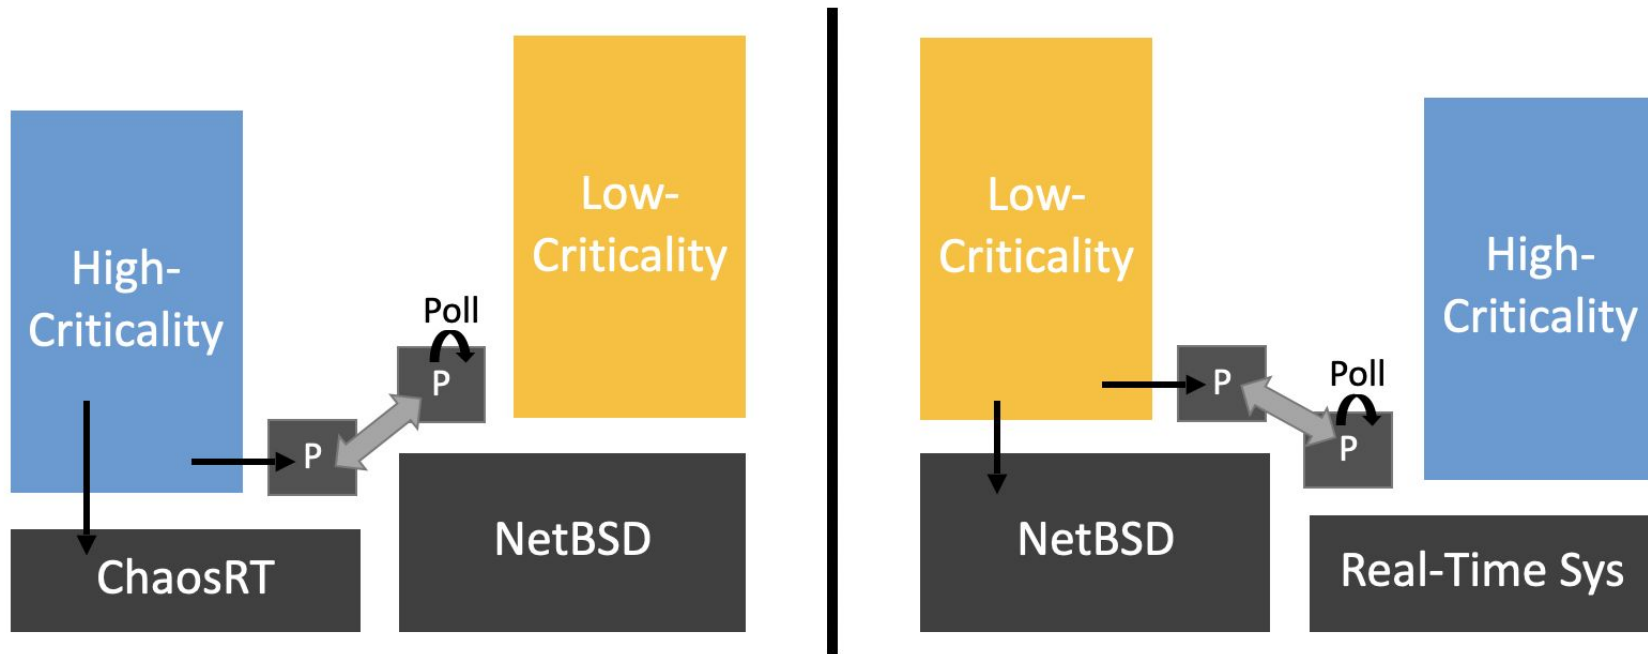
- What might happen here?
- (Recall earlier)



IPI Interference!

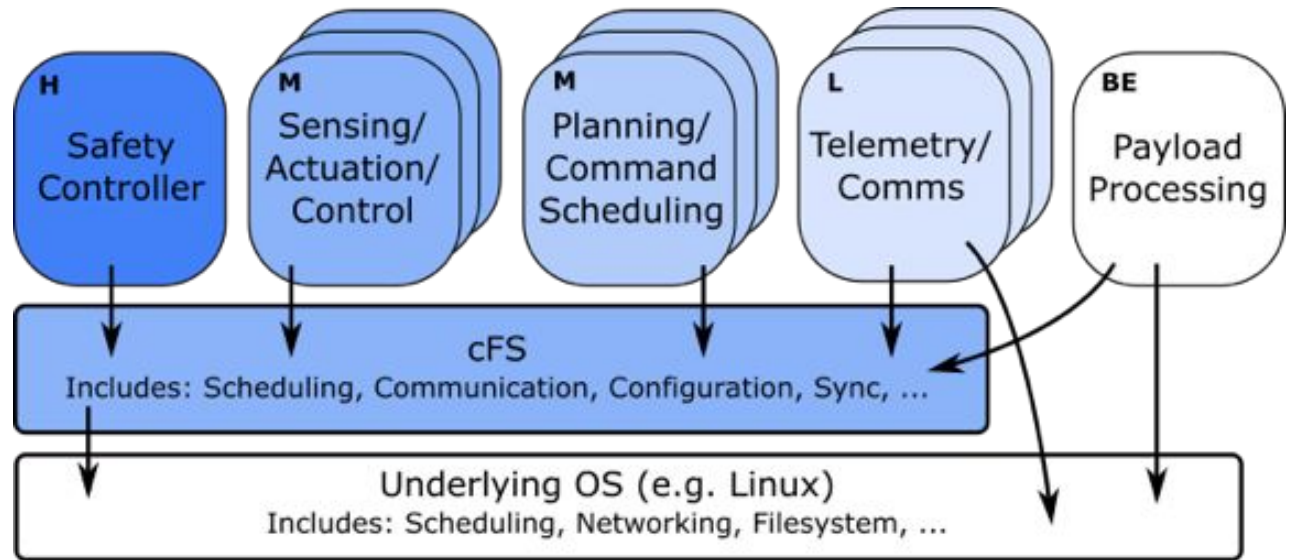



Proxies: Limiting IPI Interference



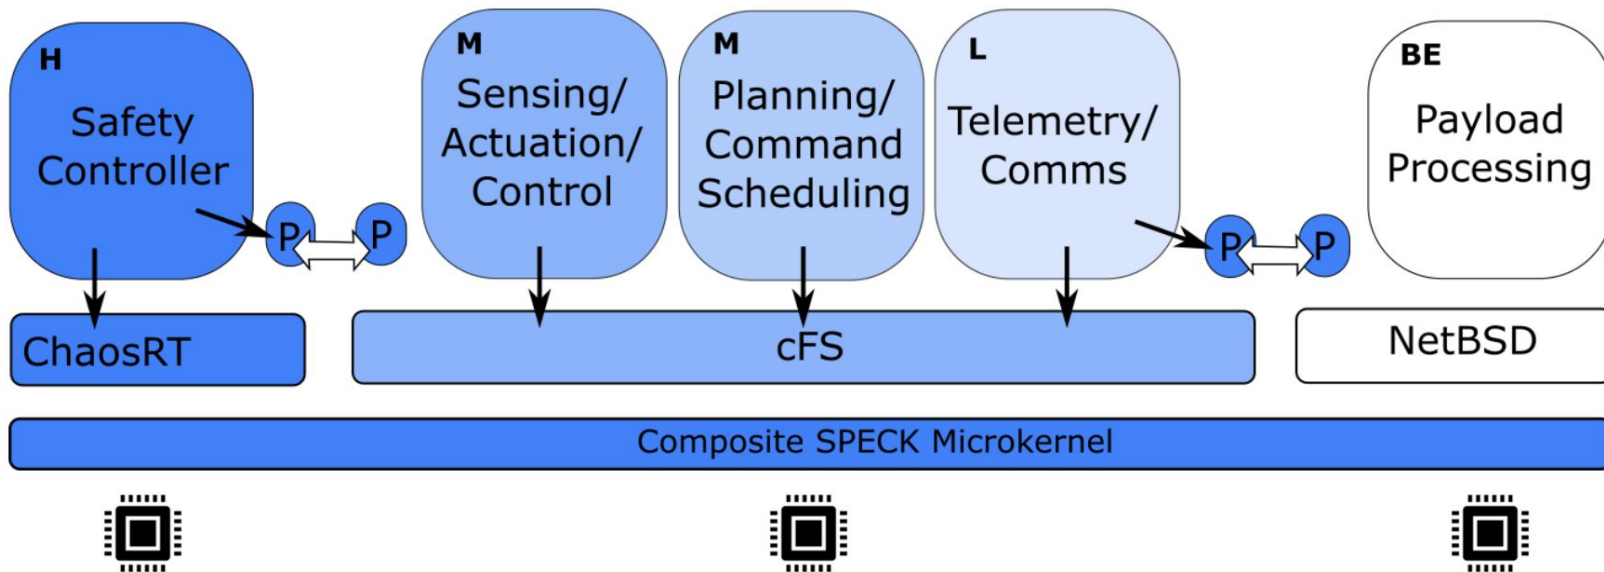
Evaluation

Use Case: The cFE (aka cFS)




cFE under CHAOS

— — —



Worst Case Communication Latency



Summary of Features

-- -- --

Problem	Composite OS	Chaos Solution
Shared-memory interference	Lockless kernel	
IPI interference	Explicit IPIs	Bound IPI interference
Configurable, isolated sched	User-level sched	
High-assurance, predictable execution		Minimal runtime
Inter-assurance-level coordination		Efficient, latency- bounded comm.
Strong inter-assurance isolation	Capability-based	Devirtualization

Critiques

Questions from the Issue

@pcodes: Could this scale to multiple assurance levels (does it even make sense to)?

@anguyen0204: How do TCAPS work?

@rebeccc: IPIs vs shared memory: why not try and resolve the issues of shared memory?

Issues from the Issue

@pcodes: The JpS (jargon per sentence) was too high

@others: The paper was too dense

@rebeccc: Devirtualization--unclear what this meant in the context of the paper

@Others: The evaluation graphs kinda suck